

Lecture 3 – Hello World & Simple Types

J. Zarnett

jzarnett@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

August 27, 2016

Acknowledgments: W.D. Bishop

Part I

Hello World

By longstanding tradition, the first program example when learning a new programming language is known as the **Hello World** program.

All it does is write the text “Hello World” to the display device.

It is a simple program; easily understood.

It takes very little time to write and is hard to get wrong.

Successful execution of it proves that the environment, compiler, display system, etc. are all present and working.

```
#include <iostream>

using namespace std;

int main()
{
    cout << "Hello World!" << endl;
    return 0;
}
```

Examining the Hello World Program

Let's examine each of the lines of this program.

At this point it is not expected that you will understand all the lines.

We'll say what they are, so more advanced students can understand the structure. If it's confusing to you, you will understand later.

What we are really interested in takes place between the inner set of { } braces.

Examining the Hello World Program

```
#include <iostream>
```

This includes a header file for input/output purposes. Including the header file makes certain functions available for use in the program.

```
using namespace std;
```

This indicates the use of functions in the std namespace.

```
int main()
```

This is the entry point for program execution.

```
cout << "Hello World!" << endl;
```

Tells the computer to output “Hello World” to the console.

It also tells the computer to end the line and move to the next one.

```
return 0;
```

Indicates this function is finished and its return value is zero.

In the Hello World program, we see the { and } braces.

These define a **block**: a grouping of statements.

The { indicates the start of a block;

The } indicates the end of that block.

They must be balanced (otherwise it's a syntax error).

Examine: `cout << "Hello World!" << endl;`

Note the presence of a semicolon (;) at the end of that statement.

Often, a single statement appears on a line.

In C++, the semicolon denotes the end of many (but not all) statements.

The end of the line doesn't indicate the end of a statement.
This allows:

- A statement to be on more than one line; or
- More than one statement on a single line.

Forgetting a semicolon is a common syntax error.

All languages provide an entry point for program execution.

The **program entry point** tells the compiler where to find the first instruction of a program.

There can be only one program entry point at compilation time.

In C++, the simplest program entry point is defined as:

```
int main( ) { }
```

Our code goes in the block between the { and } braces after `main()`.

We can, for the moment, when creating a program, ignore all the other bits and just write our code inside this “main” block.

Later on, we'll understand all the parts.

Writing & Compiling the Program

Programs are formed from source code: text written in a text file.
The text is written in a language the compiler understands (C++).

By convention, files containing C++ code use the file extension `.cpp`, such as `helloworld.cpp`

In this course, we will use the command-line compiler `gcc`.

To compile the program at the command prompt, use the command:
`gcc filename.cpp`

If we compile the Hello World program, what is our output?

On the assumption that there were no syntax errors, the compiler will run and there will be no error output.

You are left with a compiled, executable program.

If we made a mistake in writing the program, the compiler may identify it and report that to us.

[Demo: using the compiler; seeing successful/unsuccessful compiles]

If compilation did not succeed, we cannot run the program.
Because no program was produced.

The compiler provides us with some error message(s) to identify why the compile was unsuccessful.

Read these messages carefully to find the error in the program.

Compiling again without fixing the error results in the same outcome.

After we correct the error, compile the program again.

If compilation succeeded, an executable file was generated and we can run this file.

On the command line, we start execution by using the name of the executable.

By default, gcc names the output executable `a.out`.

When this program executes, it performs the following steps:

- 1 Starts executing instructions at the start of `main()`
- 2 Executes `cout << "Hello World!" << endl;` displaying the “Hello World” text on the display
- 3 Finishes execution: returns control to the operating system

[Demo: running the program]

The implementation is straightforward in a high-level language.

It's easy for us to write, but for the computer it is complex.

It requires the execution of many machine language instructions.

Each character must be written to the screen at the right time.

Output statements equal many machine language instructions.

Part II

Simple Types

Flash back to high school algebra: variables.

Algebraic variables look like this: x

We often **declare** a variable like so: Let $x = 5$.

Conceptually: x is a quantity and at the moment, the value of x is 5.

To conceptualize it visually, we have a box labelled x .
The number in that box is 5.

Variables are changeable: later, we might say: Let $x = 7$.
The number in the box changes to 7.

In algebra, you commonly saw expressions involving variables:
 $y = x + 1$

We have a new box, labelled y , and we evaluate $x + 1$.
Because $x = 7$ and $7 + 1 = 8$, $y = 8$.
Write 8 into the box labelled y .

We have the same idea in programming: software variables.

Conceptually it's the same idea: we **declare** a variable.

When a variable is declared, the computer marks a memory location to represent that variable.

That memory location is accessible by the name of the variable.

This is like creating a box and labelling it x after we write “Let $x = 5$ ”

When putting a value in the variable's box, we call that **assignment**.

Assignment can set or change the value of the variable.

Sometimes we can declare a variable to be **read-only**: we can set its value when we create it, but cannot change it after that.

In algebra, a variable x always contains a number (of some kind).

In some programming languages, when declaring a variable, it is necessary to specify a **Type**.

The type tells the compiler and computer what kind of value is stored in the box.

Example: declaring a variable as an **integer** (whole number).

Why does the type matter? Some possible answers:

Tell the compiler how much space it takes to store the value.

A big number, or one with high precision, takes more space to store.

Optimization: the compiler can translate your code into specialized instructions that operate on that data type rather than general ones.

Compile time checking: avoid assigning a “wrong” value to a variable.
e.g., prevent assigning 1.5 to a variable that should be an integer.

(Some languages don't care & just figure things out at runtime.
They do not benefit from this compile time checking.)

The equivalent to our algebraic statement “let $x = 5$ ” in C++ code:

```
int x = 5;
```

This statement follows the pattern *type identifier = value*;

`int` is the type; it's a language keyword indicating an integer variable.

`x` is the identifier; it's the name of the variable (and is case-sensitive).

`5` is a literal; it's the explicit value being stored in the variable `x`.

It's not strictly necessary to assign a value to `x` when it is declared.

```
int x;
```

is acceptable as a declaration.

At some later time, we can set the value of this variable with an assignment statement:

```
x = 3;
```

The first time a variable is assigned a value, we call that **initialization**.

It is often an error to use an uninitialized (not set) variable.

To declare a variable as read-only, we use the keyword **const**.

Example: `const int c = 0;`

If we mark a variable as `const`, it means that after assigning a value to that variable, we may not change it later.

A later attempt to assign `c = 1;` is a compile-time error.

C++ has many different variable types.

Most of them refer to different types of numbers.

There are other values we can store in variables, but we'll examine numbers first.

C++ Built-In Numeric Types

Type	Description
short int	Signed 16-bit integer
unsigned short int	Unsigned 16-bit integer
int	Signed 16-bit integer
unsigned int	Unsigned 16-bit integer
long int	Signed 32-bit integer
unsigned long int	Unsigned 32-bit integer
long long int	Signed 64-bit integer
unsigned long long int	Unsigned 64-bit integer
float	Signed 32-bit floating point
double	Signed 64-bit floating point
long double	Signed 80-bit extended precision floating point

Note that the number of bits are the minimum bits; it varies by platform.

In particular, in most cases, the regular `int` is 32 bits.

C++ Built-In Numeric Types

Type	Size in bits	Format	Value range	
			Approximate	Exact
character	8	signed (one's complement)	-127 to 127 ^[note 1]	
		signed (two's complement)	-128 to 127	
		unsigned	0 to 255	
	16	unsigned	0 to 65535	
	32	unsigned	0 to 1114111 (0x10ffff)	
integral	16	signed (one's complement)	$\pm 3.27 \cdot 10^4$	-32767 to 32767
		signed (two's complement)		-32768 to 32767
		unsigned	0 to $6.55 \cdot 10^4$	0 to 65535
	32	signed (one's complement)	$\pm 2.14 \cdot 10^9$	-2,147,483,647 to 2,147,483,647
		signed (two's complement)		-2,147,483,648 to 2,147,483,647
		unsigned	0 to $4.29 \cdot 10^9$	0 to 4,294,967,295
	64	signed (one's complement)	$\pm 9.22 \cdot 10^{18}$	-9,223,372,036,854,775,807 to 9,223,372,036,854,775,807
		signed (two's complement)		-9,223,372,036,854,775,808 to 9,223,372,036,854,775,807
		unsigned	0 to $1.84 \cdot 10^{19}$	0 to 18,446,744,073,709,551,615
floating point	32	IEEE-754 ↗	$\pm 3.4 \cdot 10^{\pm 38}$ (~7 digits)	<ul style="list-style-type: none"> min subnormal: $\pm 1.401,298,4 \cdot 10^{-47}$ min normal: $\pm 1.175,494,3 \cdot 10^{-38}$ max: $\pm 3.402,823,4 \cdot 10^{38}$
	64	IEEE-754	$\pm 1.7 \cdot 10^{\pm 308}$ (~15 digits)	<ul style="list-style-type: none"> min subnormal: $\pm 4.940,656,458,412 \cdot 10^{-324}$ min normal: $\pm 2.225,073,858,507,201,4 \cdot 10^{-308}$ max: $\pm 1.797,693,134,862,315,7 \cdot 10^{308}$

When assigning the numeric types in code, we often use literals like 5.

Literals may need a letter (or two) after them to tell the compiler the format of the literal.

If you don't specify, the compiler will try to guess what you meant.

- Unsigned numbers are followed by U or u
Examples: 5U, 5u, etc.
- Long numbers are followed by L or l
Examples: 5L, 5l, 362.f12L etc.
- Single-precision, floating point numbers are followed by F or f
Examples: 362.5F, 362.5f, etc.

We can also express numbers with powers of 10.

The double-precision floating point value for 3.625×10^2 can be expressed as 3.625E2D or 3.625e2d.

Assignment statements would therefore look like:

```
long example1 = 9001;
```

```
float example2 = 365.25f;
```

```
double example3 = 75.5e4;
```


Too Many Number Types!

There are a lot of choices, but it turns out in practice you are likely to use just a small number of them.

Typically `int` and `double` are the most common choices for integer and real (floating point) numbers.

Other types have their applications, however.

Aside from numbers, there are 2 other simple types to examine today.

The `bool` (boolean variable); and
The `char` (character).

A `bool` defines a Boolean variable.

A Boolean variable can only have one of two possible values:
`true` or `false`.

In C++, Boolean variables may be assigned values of 0 or 1.
But this is not recommended.

Example: `bool printingEnabled = true;`

A `char` defines a character variable.

A character is a single symbol from an alphabet, like 'a', 'z', or a Japanese kanji.

A `char` literal is enclosed in single quotation characters ('').

Some examples of valid `char` declarations:

`char myChar1 = 'X';` – Character Literal

`char myChar2 = '\x0058';` – Hexadecimal

`char myChar3 = '\u0058';` – Unicode

A character can also be used to represent a numeric quantity.

```
#include <iostream>

using namespace std;

int main()
{
    char myChar1 = 'X';
    char myChar2 = '\x0058';
    char myChar3 = '\u0058';

    cout << myChar1 << myChar2 << myChar3 << endl;
    return 0;
}
```

[In-Class Demo: the output of this program]

This is the second time (after the Hello World) program that we've seen console output (cout statements).

In the next lecture, we'll look at how console input and output work.